

Structured Query Language (SQL)

Lecturer: Manar Ayman

Made by: Shahinaz S. Azab

Mail: mafathi@iti.gov.eg

Edited by: Mona Saleh , Rana Salah, Manar Ayman

Room: 3005

Structured Query Language (SQL)

Data Definition
Language (DDL)

Data
Manipulation
Language (DML)

Data Control
Language (DCL)



Database Transaction

- A transaction is an executing program that forms a logical unit of database actions.
- It includes one or more database access operations such as insert, delete and update.
- The database operations that form a transaction can either be embedded within an application program or they can be specified interactively via a high-level query language such as SQL.

Database Transaction

Properties

ACID Properties

1. Atomicity

A transaction is atomic, meaning it's treated as a single, indivisible unit. It's either fully completed or fully rolled back if any part fails.

2. Consistency

Transactions bring the database from one consistent state to another. Data must satisfy predefined integrity constraints.

3. Isolation

Transactions are isolated from each other, ensuring that concurrent transactions do not interfere with each other.

4. Durability


Once a transaction is committed, its changes are permanent and survive system failures.

Database Schema

A **schema** is a group of related objects in a database.

There is one owner of a schema who has access to manipulate the structure of any object in the schema.

A schema does not represent a person, although the schema is associated with a user that resides in the database.



Data Types


A data type determine the type of the data that can be stored in a database column. The most Commonly used data types are:

Boolean

Date and Time

Alphanumeric

Numeric



Data types used to store characters, numbers, special characters, or nearly any combination.

15 8 3 10

1010101010
0101010101
1010101010
0101010101
1010101010



Database Constraints

- Primary Key (Not Null + Unique)
- Not Null
- Unique Key
- Referential Integrity (FK)
- Check



Data Definition Language (DDL)

- **Creating Database Objects:**
 - Like tables, Views, Indexes, and schemas
- **Modifying Database Objects:**
 - Such as Altering table structures or adding constraints.
- **Deleting Database objects:**
 - Delete unwanted database objects using (Drop or Truncate).

Data Definition Language (DDL) (Cont'd)

- **Data Integrity:**
 - DDL helps enforce data integrity by defining constraints such as primary keys, foreign keys, and unique constraints.
- **Data Dictionary:**
 - DDL statements are typically stored in a data dictionary, which is a repository of metadata about the database.
- **Security Implications:**
 - DDL commands are powerful and should be used with caution, as they can impact the entire database structure.

Create Command


- **Syntax:**

```
CREATE TABLE table_name  
(column1 DATA_TYPE [CONS_TYPE CONS_NAME],  
 column2 DATA_TYPE [CONS_TYPE CONS_NAME],....)
```

- **Example:**

```
CREATE TABLE Students  
(  
    ID NUMBER(15) PRIMARY KEY,  
    First_name CHAR(50) NOT NULL,  
    Last_name CHAR(50),  
    City CHAR(50),  
    Birth_Date DATE  
);
```

Apply on Create Command



Alter Command

- **Syntax:**

```
ALTER TABLE table_name  
ADD column_name DATA_TYPE;  
/ DROP COLUMN column_name;
```

- **Example:**

- ALTER TABLE Students

- ADD CONSTRAINT Found_LName NOT NULL (Last_name);

- ALTER TABLE Students
ADD country CHAR(50);

- ALTER TABLE Students
DROP COLUMN City;

Apply on ALter Command

- Add Department location as new column on Department table.
- Add constraint on Department location to be not null.
- Add Age (new column) at Dependent table.
- then delete the Age column.

Truncate Command

- **Syntax:**
 - TRUNCATE TABLE table_name;
- **Example:**
 - TRUNCATE TABLE Students;

Drop Command

- **Syntax:**
 - DROP TABLE table_name;
- **Example:**
 - DROP TABLE Students;

- **Truncate** used to delete the data from the table.
- **Drop** used to delete the table structure with its data.

Data Manipulation Language (DML)

- **Retrieving Data:**
 - Retrieve data from one or more database tables using queries (select Command).
- **Inserting Data:**
 - Insert new records or rows into database tables (Insert Command).
- **Updateing Data:**
 - Updating existing data, making it useful for modifying records (Update Command).
- **Deleting Data:**
 - Delete records or rows from database tables (Delete Command).

Data Manipulation Language (DML)

- **Data Filtering:**
 - DML allows you to filter data using conditions specified in the queries (e.g., WHERE clause).
- **Data Transformation:**
 - DML supports data transformation through functions and calculations within queries.
- **Data Integrity:**
 - DML statements should adhere to data integrity constraints defined using DDL (Data Definition Language).

Insert Command

- **Syntax:**
 - a. **Specify both the column name and the values to be inserted:**

```
INSERT INTO table_name  
(column1, column2, column3, ....)  
VALUES ( value1 , value2, value3,.....);
```

- **Example:**

```
INSERT INTO Customers  
(Customer_Name, Contact_Name, Address, City,  
Postal_Code, Country)  
VALUES ('Cardinal', 'Tom B. Erichsen', 'Skagen 21',  
'Stavanger', '4006', 'Norway');
```

Insert Command

- **Syntax:**
 - b. If you are adding values for all the columns of the table, you do not need to specify the column names in the SQL query:

```
INSERT INTO table_name  
VALUES (value1, value2, value3, ...);
```
- **Example:**

```
INSERT INTO Customers  
VALUES ('Cardinal', 'Tom B. Erichsen', 'Skagen 21',  
'Stavanger', '4006', 'Norway');
```

Insert Command

- **Syntax:**
 - c. **Insert multiple rows:**

- **Example:**
 - `INSERT INTO Customers (CustomerName,
ContactName, Address, City, PostalCode, Country)
VALUES
('Cardinal', 'Tom B. Erichsen', 'Skagen 21', 'Stavanger',
'4006', 'Norway'),
('Greasy Burger', 'Per Olsen', 'Gateveien 15', 'Sandvika',
'4306', 'Norway');`

Update Command

- **Syntax:**
 - UPDATE table_name
SET column1= value1, column2= value2, ...
WHERE condition;
- **Example:**
 - **Update Single record:**
 - UPDATE Customers
 - SET ContactName = 'Alfred Schmidt', City= 'Frankfurt'
 - WHERE CustomerID = 1;

Update Command

- Update multiple record:
 - UPDATE Customers
 - SET ContactName='Juan'
 - WHERE Country='Mexico';
- **Update Warning:**
 - Be carful when updating records. If you omit the where clause, All records will be updated!
- **Example:**
 - UPDATE Customers
 - SET ContactName='Juan';

Delete Command

- **Syntax:**
 - DELETE FROM table_name WHERE condition;
- **Example:**
 - DELETE FROM Customers WHERE CustomerName='Alfreds Futterkiste';
- **Delete all records:**
 - DELETE FROM table_name;
- **Example:**
 - DELETE FROM Customers;

Truncate vs Delete


Truncate	Delete
Delete the content of table	Delete the content of table
It's a DDL Command	It's a DML Command
Can't use where clause	Can use where clause to delete specific records
The data page is locked before removing the table data	A tuple is locked before removing it

Truncate vs Delete

Truncate	Delete
Faster than delete	Slower than delete
Need alter permission on the table	need delete permission on the table
Can't be used with indexed view	can be used with indexed view
Can't active trigger	can active trigger
occupies less transaction spaces than delete	occupies more transaction spaces than truncate

Simple Queries

- **Syntax:**
 - SELECT <attribute list>
FROM <table list>
WHERE <condition>
ORDER BY <attribute list>;




Simple Queries

- **Select all data:**
 - `SELECT *
FROM departments;`
- **Select specific data:**
 - `SELECT emp_id , emp_name, dept_id
FROM employees
WHERE location = "Cairo";`

Distinct Keyword

- It's a row keyword that displays unique rows:
- Example:
 - SELECT DISTINCT DNO
FROM employees;

EmpNO	Name	DNO	JobID
100	Ahmed	2	Sales_Rep
200	Mai	2	It_PROG
300	Ali	3	Sales_Rep



Distinct Keyword

- It's a row keyword that displays unique rows:
- Example:
 - SELECT DISTINCT DNO,JobID FROM employees;

EmpNO	Name	DNO	JobID
100	Ahmed	2	Sales_Rep
200	Mai	2	It_PROG
300	Ali	3	Sales_Rep

DNO	JobID
2	Sales_Rep
2	IT_Prog
3	Sales_Rep

Comparison & Logical Operators:

- = , > , < -----> **Single Row Operator:**
 - WHERE Salary >= 1500 and Salary <> 2500;
 - WHERE Super_SSN = 321 OR Super_SSN = 321;
- **ALL, ANY, IN** -----> **Multi-row Operator:**
 - WHERE Super_SSN IN (321 , 223);
- **AND , OR** ----->**Combained more than 1 condition**
 - WHERE Salary between 1500 and 2500;
- **Like** -----> **Compare based on pattern**
 - WHERE fname like '?o*';
 - ? -----> 1 char
 - * -----> Zero or more char

Arithmetic Expressions

- ```
SELECT last_name , Salary , Salary+300
 FROM Employees;
```
- Order of Precedence: \*, /, +, -
  - You can enforce priority by adding parentheses
- ```
SELECT last_name , Salary  
,10* (Salary+300) as proposed salary  
      FROM Employees;
```

$$\frac{x^a}{x^b} = x^{a-b}$$



Order by Clause

Used to sort the result-set in Ascending or Descending order (Asc order is the default)

- **Syntax:**
 - `SELECT column1, column2, ...
FROM table_name
ORDER BY column1, column2, ... ASC|DESC;`
- **Example:**
 - `SELECT * FROM Products
ORDER BY Price;`

Types of Join

- **INNER JOIN:**
 - Returns rows when there is at least one match in both tables.
- **OUTER JOIN:**
 - Returns matched rows as well as rows when there is no match in one of the tables.
- **CARTESIAN PRODUCT:**
 - Returns all possible combination of rows.



Apply on Join

- Retrieve the name, address of all employees who work for Research Department

```
SELECT fname, Lname, Address  
FROM Employee , Department  
WHERE Dname= “research” and  
Department.number = employee.Dno;
```

Table Alias:

- You can add alternative name (Alias) for Table name.
- Alias can be useful for long or complex table names.
- You can use table alias instead of table name to resolve ambiguity.

```
SELECT e.ID , d.name , e.name , e.Salary  
FROM department d , employee e  
WHERE d.id = e.deptid  
ORDER BY d.name;
```

Self Join:

- Is a join in which a table is joined with itself (Unary relationship), Specially when the table has a foreign key which referenced its own PK.
- To Join a table itself means that each row of the table is combined with itself and with every other row of the table.
- The self-join can be viewed as a join of 2 copies of the same table.
- Example:
 - `SELECT e.name Employee_Name , s.name Supervisor
FROM employee e , employee s
WHERE e.supervisorID = s.ID;`

Outer Join

Left Join

Return all rows from the left table, even if there are no matches in the right table.

Right Join

Return all rows from the right table, even if there are no matches in the left table.

Full Outer

Returns rows when there is a match in one of the tables.

Apply on Outer Join:

- Display all departments' information even if they have no employees assigned.
- ```
SELECT e.name AS Employee,
d.dept_id, d.name AS Department
FROM employees e RIGHT
OUTER JOIN departments d
ON e.dept_id = d.id;
```

| Emp_name   | Dept ID | Dept      |
|------------|---------|-----------|
| Ahmed Ali  | 100     | IT        |
| Amr Samir  | 200     | Marketing |
| Mona Selim | 100     | IT        |
|            | 300     | HR        |

# Equi and Non-Equi Join

- **Equi join:** when the join condition is based on the equality operator (Ex. D.dno=e.deptno)
- **Non-Equi join:** when the join condition is based on any operator rather than the equality operator

# Non-Equi Join Example



- Display employees information ( name , salary , title) with salary grade for each employee.
- ```
SELECT e.name, e.salary, j.grade
FROM employees e, job_grades j
WHERE e.salary BETWEEN j.lowsal
AND j.highsal;
```

LowSal	HighSal	Grade
1000	4000	B
4000	7000	A



Sub Queries

- Sub-Query (Nested Query): is a complete SELECT query inside another SELECT
- The Inner query is executed first then the outer query
- The inner query is usually placed in the WHERE or HAVING clauses
- Sometimes it is placed in the FROM clause and called “inline view”

Example

- Find the names of employees whose working locations are Giza

```
SELECT      name
FROM        Employee
WHERE       Dno    IN  (
                      SELECT Dnumber
                      FROM  Dept
                      WHERE Location = "Giza"
                );
```

Example (Cont'd)

- Display department name for the highest paid employee

```
SELECT      dname
FROM        Dept
WHERE       deptno = (
              SELECT deptno
              FROM   emp
              WHERE sal = (
                            SELECT MAX (sal)
                            FROM emp
                          )
            );

```

Example (Cont'd)

- Find the names of employees whose salary is greater than the salary of the employees in department 5

```
SELECT      Lname , Fname  
FROM        employee  
WHERE       salary > ALL  (  
                      SELECT salary  
                      FROM   employee  
                      WHERE Dno = 5  
                     );
```

Union Operator:

- UNION operator is one of the set operators that unifies the output of two select statements into one result.
- Two conditions must apply for the two selects:
 - Same number of columns
 - Same data type of columns
- The result displays the names of the first query.
- If one row is contained in the result of the two selects it is displayed only once.
- UNION ALL is functionally equivalent to UNION operator, except that it doesn't eliminate repetition in the result set.

Union Operator Example

- Find the departments numbers whose location in GIZA or manager no = 10

```
SELECT Dnumber  
FROM department  
WHERE MGRSSN = 10;
```

UNION

```
SELECT Dnumber  
FROM dept_locations  
WHERE dlocation = "GIZA";
```



Union Operator Example

- Display names of current employees as well as previous employees

```
SELECT Name  
FROM Employee ;
```

UNION

```
SELECT Name  
FROM Employees_retired;
```



Name
Ahmed
Mona
Samir

Correlated Sub-Query:

- When the inner query references one of the attributes of the outer query.
- The inner query is executed once per row of the outer query
- **Example:**
 - ```
SELECT product_name, list_price, category_id
 FROM production.products p1
 WHERE list_price IN (SELECT MAX (p2.list_price)
 FROM production.products p2
 WHERE p2.category_id = p1.category_id
 GROUP BY p2.category_id)
 ORDER BY category_id, product_name;
```

# Exists Keyword

- The EXISTS keyword is used with correlated sub-queries
- The EXISTS condition is considered "to be met" if the sub-query returns at least one row.
- **Syntax**
  - SELECT columns  
FROM tables  
WHERE EXISTS (sub-query );

# Exists Keyword Example

- Display suppliers' information who has orders.

- SELECT \*

```
FROM suppliers
```

```
WHERE EXISTS (
```

```
 SELECT *
```

```
 FROM orders
```

```
 WHERE
```

```
 suppliers.supplier_id = orders.supplier_id
```

```
);
```

# Exists Keyword Example

- Retrieve the name of employees who have no dependents

- SELECT name

```
FROM employee
```

```
WHERE NOT EXISTS (
```

```
 SELECT *
```

```
 FROM dependent
```

```
 WHERE ssn=Essn
```

```
);
```

# Exists Condition with DML

- DELETE FROM suppliers

WHERE NOT EXISTS (

SELECT \*

FROM orders

WHERE

    suppliers.supplier\_id = orders.supplier\_id

);

# Aggregate Functions:

- Aggregate Functions ( group functions): perform a specific operation on a number of rows and return one result per group
- Example:
  - COUNT, SUM, MAX, MIN, and AVG
- Group Functions ignore null values in the columns.

# Aggregate Functions Example:

- Find the sum, maximum, minimum and average salaries of all employees.

```
SELECT SUM (salary), MAX (salary), MIN (salary), AVG (salary)
FROM Employees;
```

- Find the total number of employees in the company?

```
SELECT COUNT(*)
FROM employees;
```

- Find the number of employees in the research department?

```
SELECT COUNT(*)
FROM employee , department
WHERE dno=dnumber AND dname ='Research';
```

# Grouping

- If you want to apply aggregate functions to subgroups of tuples, use GROUP BY clause
- GROUP BY clause must be used in conjunction with aggregate functions
- You can filter group results using HAVING clause
- HAVING clause is used for applying conditions on group functions

# Grouping Examples

- For each department, retrieve the department number, the number of employees in the department, and their average salaries.

```
SELECT dno , COUNT (*)
, AVG (salary)
FROM employee
GROUP BY dno
```

| Employee | Dno | Salary | Dno | Salary |
|----------|-----|--------|-----|--------|
| Ahmed    | 1   | 1000   | 1   | 3000   |
| Mai      | 2   |        | 2   | 2000   |
| Mohamed  | 1   | 5000   |     |        |
| Hosam    | 2   | 2000   |     |        |

# Grouping Examples

- For each project on which more than two employees work, retrieve the project number, the project name , and the number of employees who work on the project

```
SELECT pnumber, pname ,count (works_on.pno)
FROM project , Works_on
WHERE Pnumber = Pno
GROUP BY pnumber, pname
HAVING COUNT (*) > 2
ORDER BY Pnumber;
```

# Data Control Language

- **Access Control:**
  - DCL is primarily concerned with managing who can access specific data or perform certain actions within a database.
  - It defines and enforces security and authorization rules.
- **Two Main DCL Statements:**
  - There are two main DCL statements in SQL: GRANT and REVOKE.
  - These statements are used to grant or revoke specific privileges to users or roles.

# Data Control Language

Grant


- GRANT SELECT ON TABLE employees TO Ahmed;
- GRANT ALL ON TABLE department TO Mary,  
Ahmed;
- GRANT SELECT ON TABLE employees TO Ahmed  
WITH GRANT OPTION;

Revoke

- REVOKE UPDATE ON TABLE department  
FROM Mary;
- REVOKE ALL ON TABLE department FROM Mary,  
Ahmed;


# Views:

- A view is a logical table based on a table or another view
- A view contains no data of its own, but is like a window through which data from tables can be viewed or changed
- The tables on which a view is based are called base tables
- The view is stored as a SELECT statement in the data dictionary.



# Advantages of Views:

- Restrict data access
- Make complex queries easy
- Provide data independence
- Present different views of the same data




# Simple Views and Complex Views:

| Feature                       | Simple View | Complex View |
|-------------------------------|-------------|--------------|
| Number of Tables              | One         | One or more  |
| Contain Functions             | ✗           | ✓            |
| Contain groups of data        | ✗           | ✓            |
| DML operations through a view | ✓           | Not always   |



# Creating Views:


- CREATE VIEW view [ (column 1 [ , column2 ]  
... ) ]  
AS subquery  
[ With Check Option ];



# Creating Views (Cont'd):


- Create a view to display employee names and total hours employee worked on a project

```
CREATE VIEW vw_work_hrs
AS
SELECT Fname , Lname , Pname , Hours
FROM Employee, Project , Works_on
WHERE SSN=ESSN AND PNO=PNUMBER;
```




# Views with Check Option:

- CREATE VIEW Suppliers  
AS  
SELECT \*  
FROM suppliers  
WHERE status > 15  
WITH CHECK OPTION;




# Modifying a View:

- **Syntax:**
  - CREATE OR REPLACE VIEW view\_name  
AS  
Sub-query
- **Example:**
  - CREATE OR REPLACE VIEW vw\_work\_hrs  
AS  
SELECT Fname , Lname , Pname , Hours  
FROM Employee, Project , Works\_on  
WHERE SSN=ESSN AND PNO=PNUMBER AND Dno = 5;



# Removing a View:

- **Syntax:**
  - `DROP VIEW view_name;`
- **Example:**
  - `DROP VIEW vw_work_hrs`




# Indexes:

- They are used to speed up the retrieval of records in response to certain search conditions.
- May be defined on multiple columns
- Can be created by the user or by the DBMS
- Are used and maintained by the DBMS





# Indexes (Cont'd):



File (index)

Supplier file (data)



# Indexes (Cont'd):

When to create an index:

|   |                                                                                                          |
|---|----------------------------------------------------------------------------------------------------------|
| ✓ | A column contains a wide range of values                                                                 |
| ✓ | A column contains a large number of null values                                                          |
| ✓ | One or more columns are frequently used together in a WHERE clause or a join condition                   |
| ✓ | The table is large and most queries are expected to retrieve less than 2% to 4% of the rows in the table |



# Indexes (Cont'd):

## When to don't create an index

- ✗ The columns are not often used as a condition in the query
- ✗ The table is small or most queries are expected to retrieve more than 2% to 4% of the rows in the table
- ✗ The table is updated frequently
- ✗ The indexed columns are referenced as part of an expression -- [Upper(Lname)]



# Indexes (Cont'd):

- **Creation**
  - `CREATE INDEX index_name ON Table_name (column_name);`
  - `CREATE INDEX emp_inx ON Employee (Salary);`
- **Removing**
  - `DROP INDEX index_name;`
  - `DROP INDEX emp_inx;`

# SQL Tutorials

- **Oracle SQL URL:**
  - [beginner-sql-tutorial.com](http://beginner-sql-tutorial.com)
- **Sybase SQL URL:**
  - <http://infocenter.sybase.com/help/index.jsp>
- **ANSI SQL URL:**
  - [http://www.w3schools.com/SQL/sql\\_join.asp](http://www.w3schools.com/SQL/sql_join.asp)
- **MS SQL URL:**
  - <http://msdn.microsoft.com/en-us/library/bb264565.aspx>
- **IBM Informix SQL:**
  - <http://publib.boulder.ibm.com/infocenter/idshelp/v10/index.jsp?topic=/com.ibm.sqlt.doc/sqltmst104.htm>

# Questions?

